

```
In [9]: import numpy as np
import pandas as pd

In [10]: data=pd.read_csv("customer_churn.csv")

In [11]: data
```

Out[11]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceP |
|------|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|---------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | No | DSL | Yes | ... | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | No phone service | DSL | Yes | ... | |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | No | Fiber optic | No | ... | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 7038 | 6840-RESVB | Male | 0 | Yes | Yes | 24 | Yes | Yes | DSL | Yes | ... | |
| 7039 | 2234-XADUH | Female | 0 | Yes | Yes | 72 | Yes | Yes | Fiber optic | No | ... | |
| 7040 | 4801-JAZL | Female | 0 | Yes | Yes | 11 | No | No phone service | DSL | Yes | ... | |
| 7041 | 8361-LTMKD | Male | 1 | Yes | No | 4 | Yes | Yes | Fiber optic | No | ... | |
| 7042 | 3186-AJIEK | Male | 0 | No | No | 66 | Yes | No | Fiber optic | Yes | ... | |

7043 rows × 21 columns

```
In [12]: data.head()
```

Out[12]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProte |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | No | DSL | Yes | ... | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | No phone service | DSL | Yes | ... | |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | No | Fiber optic | No | ... | |

5 rows × 21 columns

A)Data Manipulation

a) Extract the 5th column and store it in 'customer_5'

```
In [13]: customer_5=data.iloc[:,4]
customer_5
```

Out[13]:

```
0      No
1      No
2      No
3      No
4      No
...
7038   Yes
7039   Yes
7040   Yes
```

```

7041      No
7042      No
Name: Dependents, Length: 7043, dtype: object

```

b)extract the 15th column and store it in 'customer_15'

```

In [14]: customer_15=data.iloc[:,14]
customer_15

```

```

Out[14]: 0      No
1      No
2      No
3      No
4      No
...
7038   Yes
7039   Yes
7040   No
7041   No
7042   Yes
Name: StreamingMovies, Length: 7043, dtype: object

```

```

In [15]: data.head(2)

```

```

Out[15]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  MultipleLines  InternetService  OnlineSecurity  ...  DeviceProte
0      7590-  VHVEG   Female             0     Yes           No         1           No      No phone service      DSL              No  ...
1      5575-  GNVDE    Male             0     No           No        34          Yes           No      DSL              Yes  ...

```

2 rows × 21 columns

C)Extract all the male senior citizen whose payment method is electronic check and store the result in 'senior_male_electronic'

```

In [17]: senior_male_electronic=data[(data['gender']=='Male') & (data['SeniorCitizen']==1) & (data['PaymentMethod']=='Elec

```

```

In [18]: senior_male_electronic

```

```

Out[18]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  MultipleLines  InternetService  OnlineSecurity  ...  DeviceProte
20      8779-  QRDMV    Male             1     No           No         1           No      No phone service      DSL              No  ...
55      1658-  BYGOY    Male             1     No           No        18          Yes           Yes      Fiber optic      No  ...
57      5067-  XJQFU    Male             1     Yes          Yes        66          Yes           Yes      Fiber optic      No  ...
78      0191-  ZHSKZ    Male             1     No           No        30          Yes           No      DSL              Yes  ...
91      2424-  WVHPL    Male             1     No           No         1          Yes           No      Fiber optic      No  ...
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
6837     6229-  LSCKB    Male             1     No           No         6          Yes           No      Fiber optic      No  ...
6894     1400-  MMYXY    Male             1     Yes          No         3          Yes           Yes      Fiber optic      No  ...
6914     7142-  HVGBG    Male             1     Yes          No        43          Yes           Yes      Fiber optic      No  ...
6967     8739-  WWKDU    Male             1     No           No        25          Yes           Yes      Fiber optic      No  ...
7032     6894-  LFHLY    Male             1     No           No         1          Yes           Yes      Fiber optic      No  ...

```

298 rows × 21 columns

D)Extract all those customers whose tenure is greater than 70 months or their Monthly charges is more than 100 & store the result in 'customer_total_tenure'

In [19]: `data.head(2)`

Out[19]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProte |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |

2 rows × 21 columns

In [20]: `customer_total_tenure=data[((data['tenure']>70) | (data['MonthlyCharges'] >100))]`

In [21]: `customer_total_tenure`

Out[21]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DevicePro |
|------|------------|--------|---------------|---------|------------|--------|--------------|---------------|-----------------|---------------------|-----|-----------|
| 8 | 7892-POOKP | Female | 0 | Yes | No | 28 | Yes | Yes | Fiber optic | No | ... | |
| 12 | 8091-TTVAX | Male | 0 | Yes | No | 58 | Yes | Yes | Fiber optic | No | ... | |
| 13 | 0280-XJGEX | Male | 0 | No | No | 49 | Yes | Yes | Fiber optic | No | ... | |
| 14 | 5129-JLPIS | Male | 0 | No | No | 25 | Yes | No | Fiber optic | Yes | ... | |
| 15 | 3655-SNQYZ | Female | 0 | Yes | Yes | 69 | Yes | Yes | Fiber optic | Yes | ... | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 7023 | 1035-IPQPU | Female | 1 | Yes | No | 63 | Yes | Yes | Fiber optic | No | ... | |
| 7034 | 0639-TSIQW | Female | 0 | No | No | 67 | Yes | Yes | Fiber optic | Yes | ... | |
| 7037 | 2569-WGERO | Female | 0 | No | No | 72 | Yes | No | No | No internet service | ... | N |
| 7039 | 2234-XADUH | Female | 0 | Yes | Yes | 72 | Yes | Yes | Fiber optic | No | ... | |
| 7042 | 3186-AJIEK | Male | 0 | No | No | 66 | Yes | No | Fiber optic | Yes | ... | |

1259 rows × 21 columns

E)Extract all the customers whose contract is of two years,payment method is mailed check & the value of churn is 'yes' & store the result in'two_mail_yes'

In [22]: `data.head(2)`

Out[22]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProte |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |

2 rows × 21 columns

In [24]: `two_mail_yes=data[((data['Contract']=='Two year') & (data['PaymentMethod']=='Mailed check') & (data['Churn']=='Yes'))]`

In [25]: `two_mail_yes`

Out[25]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceP |
|------|------------|--------|---------------|---------|------------|--------|--------------|---------------|-----------------|---------------------|-----|---------|
| 268 | 6323-AYBRX | Male | 0 | No | No | 59 | Yes | No | No | No internet service | ... | N |
| 5947 | 7951-QKZPL | Female | 0 | Yes | Yes | 33 | Yes | Yes | No | No internet service | ... | N |
| 6680 | 9412-ARGBX | Female | 0 | No | Yes | 48 | Yes | No | Fiber optic | No | ... | |

3 rows × 21 columns

F) Extract 333 random records from the customer_churn dataframe & store the result in 'customer_333'

In [19]:

```
customer_333=data.sample(n=333)
```

In []:

```
customer_333
```

Out[]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceP |
|------|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|---------------------|-----|---------|
| 1589 | 7351-KYHQH | Female | 1 | No | No | 7 | Yes | No | DSL | No | ... | |
| 382 | 8204-YJCLA | Male | 1 | Yes | Yes | 72 | No | No phone service | DSL | Yes | ... | |
| 1480 | 8898-KASCD | Male | 0 | No | No | 39 | No | No phone service | DSL | No | ... | |
| 5645 | 4942-VZZOM | Male | 0 | Yes | No | 64 | Yes | Yes | DSL | Yes | ... | |
| 5379 | 6284-KMNUF | Female | 0 | Yes | No | 56 | Yes | Yes | Fiber optic | Yes | ... | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 106 | 6728-DKUCO | Female | 0 | Yes | Yes | 72 | Yes | Yes | Fiber optic | Yes | ... | |
| 6521 | 1092-WPIVQ | Female | 0 | Yes | Yes | 18 | Yes | Yes | No | No internet service | ... | N |
| 4810 | 1112-CUNAO | Female | 1 | No | No | 15 | Yes | Yes | Fiber optic | No | ... | |
| 4194 | 2961-VNFKL | Female | 0 | Yes | No | 71 | Yes | Yes | No | No internet service | ... | N |
| 6317 | 7493-TPUWZ | Male | 0 | No | No | 1 | Yes | Yes | Fiber optic | No | ... | |

333 rows × 21 columns

g) Get the count of different levels from the 'churn' column

In [26]:

```
data['Churn'].value_counts()
```

Out[26]:

```
No      5174
Yes     1869
Name: Churn, dtype: int64
```

Data visualization

A) Build a barplot for the 'InternetService' column

i) set x-axis label to 'Categories of internet Service'

ii) set y-axis label to 'Count of Categories'

iii)set the title of plot to be 'Distribution of internet Service'

iv)set the color of bars to be 'orange'

```
In [27]: data.head(2)
```

```
Out[27]:
```

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProte |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |

2 rows × 21 columns

```
In [28]: import matplotlib.pyplot as plt
```

```
In [ ]: data['InternetService'].value_counts()
```

```
Out[ ]: Fiber optic    3096
DSL                2421
No                 1526
Name: InternetService, dtype: int64
```

```
In [29]: x=data['InternetService'].value_counts().keys()
```

```
In [30]: x
```

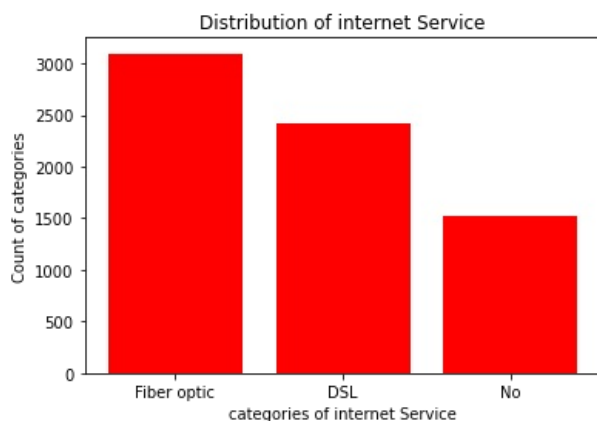
```
Out[30]: Index(['Fiber optic', 'DSL', 'No'], dtype='object')
```

```
In [32]: y=data['InternetService'].value_counts().tolist()
y
```

```
Out[32]: [3096, 2421, 1526]
```

```
In [33]: import numpy as np
import matplotlib.pyplot as plt
```

```
plt.bar(x,y,color='red')
plt.xlabel("categories of internet Service")
plt.ylabel("Count of categories")
plt.title("Distribution of internet Service")
plt.show()
```




b)Build a histogram for the 'tenure' column:

i)set num of bins to be 30

ii)set the color of bins to be 'green'

iii)Assign the title 'Distribution of tenure'

```
In [34]: plt.hist(data['tenure'],bins=30,color='green')
plt.title("Distribution of tenure")
plt.show()
```



Built a scatter-plot between 'MonthlyCharges' & 'tenure'. Map 'MonthlyCharges' to the y-axis & 'tenure' to the 'x-axis'

i)Assign the points a color of 'brown'

ii)Set the x-axis label to tenure of customer

iii)Set the y-axis label to 'Monthly Charges of customer'

iv)Set the title to the 'Tenure vs Monthly Charges'

```
In [35]: import pandas as pd
```

```
In [37]: data=pd.read_csv("customer_churn.csv")
```


```
In [ ]: data.head(2)
```


| Out[]: | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProte |
|---------|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |

2 rows × 21 columns

```
In [38]: import matplotlib.pyplot as plt
```

```
In [39]: plt.scatter(x=data['tenure'].head(20),y=data['MonthlyCharges'].head(20),color='brown')
plt.xlabel('tenure of customer')
plt.ylabel('Monthly Charges of customer')
plt.title('Tenure vs Monthly Charges')
plt.show()
```






d)built a boxplot between tenure & contract. Map 'tenure' on y-axis and 'Contract' on the x-axis

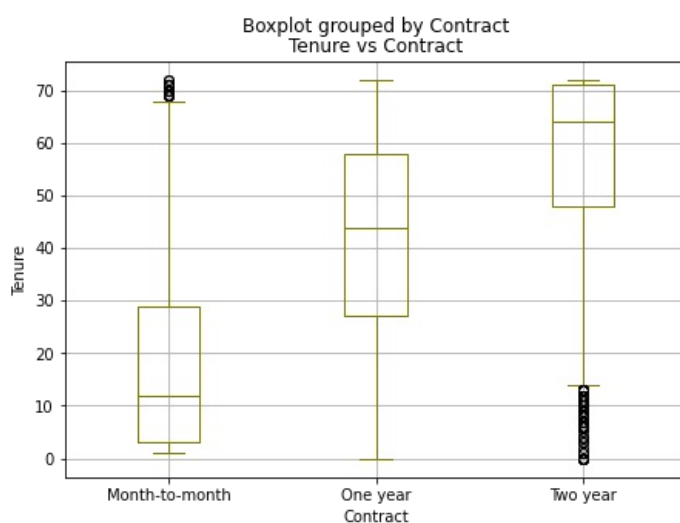
```
In [40]: import seaborn as sns
```

```
In [41]: sns.boxplot(x=data['Contract'],y=data['tenure'])
plt.show()
```



```
In [4]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [42]: data.boxplot(by=['Contract'],column="tenure",figsize=(7,5),color='olive')
plt.xlabel("Contract")
plt.ylabel("Tenure")
plt.title("Tenure vs Contract")
plt.show()
```



Linear Regression

A)Built a simple linear model where dependent variable is 'MonthlyCharges' and independent variable is 'tenure'

- i) Divide the dataset into train and test sets is 70:30 ratio
- ii) built the model on train set and predict the values on test set
- iii) After predicting the value find the root mean square error
- iv) Find out the error in prediction & store the result in 'error'
- v) find the root mean square error

In [43]: `from sklearn.model_selection import train_test_split`

In [44]: `from sklearn.linear_model import LinearRegression`

In [45]: `data.head(2)`

Out[45]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProte |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |

2 rows × 21 columns

In [46]: `x=pd.DataFrame(data['tenure']) #independent variable`

In [47]: `y=data['MonthlyCharges'] # dependent variable`

In [50]: `x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)`

In [51]:

```
print(data.shape)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(7043, 21)
(4930, 1)
(2113, 1)
(4930,)
(2113,)
```

In [52]:

```
import numpy as np
import pandas as pd
```

In [55]: `data=pd.read_csv("customer_churn.csv")`

In [56]: `data.head(2)`

Out[56]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProte |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|-------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |

2 rows × 21 columns

In [57]: `lr=LinearRegression()`

In [58]:


```
lr.fit(x_train,y_train)
```

```
Out[58]: LinearRegression()
```

```
In [60]: y_pred=lr.predict(x_test)
```

```
In [61]: y_pred
```

```
Out[61]: array([60.95089608, 72.98096699, 59.1903979 , ..., 75.62171426,
              70.63363608, 65.6455579 ])
```

```
In [62]: y_test.values
```

```
Out[62]: array([ 58.2 , 116.6 ,  71.95, ..., 109.95,  24.55,  81.6 ])
```

```
In [63]: from sklearn.metrics import mean_squared_error
import numpy as np
```

```
In [64]: msc=mean_squared_error(y_test,y_pred)
```

```
In [65]: error=np.sqrt(msc)
```

```
In [66]: error
```

```
Out[66]: 29.394584027273893
```

D)Logistic Regression

a)Built a simple logistic regressin model where dependent variable is 'churn' & independent variable is 'MonthlyCharges'

i)Divide the dataset into 65:35 ratio

ii)Build the model on train set and predict the values on test set

iv)Build the confusion matrix and get the accuracy score

```
In [67]: from sklearn.linear_model import LogisticRegression
```

```
In [68]: x=pd.DataFrame(data['MonthlyCharges'])
```

```
In [69]: y=data['Churn']
```

```
In [70]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.65,random_state=0)
```

```
In [71]: logmodel=LogisticRegression()
```

```
In [72]: logmodel.fit(x_train,y_train)
```

```
Out[72]: LogisticRegression()
```

```
In [73]: y_pred=logmodel.predict(x_test)
```

```
In [74]: y_pred
```

```
Out[74]: array(['No', 'No', 'No', ..., 'No', 'No', 'No'], dtype=object)
```

```
In [75]: y_test.values
```

```
Out[75]: array(['No', 'No', 'No', ..., 'Yes', 'No', 'No'], dtype=object)
```

```
In [77]: from sklearn.metrics import confusion_matrix, accuracy_score
```

```
In [78]: confusion_matrix(y_pred, y_test)
```

```
Out[78]: array([[1815, 651],
               [ 0, 0]])
```

```
In [79]: (1815+0)/(1815+651+0+0)
```

```
Out[79]: 0.7360097323600974
```

```
In [80]: accuracy_score(y_pred, y_test)
```

```
Out[80]: 0.7360097323600974
```

b) Built a multiple logistic regression model where dependent variable is 'Churn' & independent variable are 'tenure' & 'MonthlyCharges'

i) Divide the dataset in 80:20 ratio

ii) Built the model on train set and predict the values on test set

iii) Built the confusion matrix and get accuracy score

```
In [81]: x=pd.DataFrame(data.loc[:,['tenure','MonthlyCharges']])
```

```
In [82]: y=data['Churn']
```

```
In [83]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.8,random_state=0)
```

```
In [84]: mlr=LogisticRegression()
```

```
In [86]: mlr.fit(x_train,y_train)
```

```
Out[86]: LogisticRegression()
```

```
In [87]: y_pred=mlr.predict(x_test)
```

```
In [88]: y_pred
```

```
Out[88]: array(['No', 'No', 'No', ..., 'No', 'No', 'No'], dtype=object)
```

```
In [89]: y_test.values
```

```
Out[89]: array(['No', 'No', 'No', ..., 'Yes', 'No', 'No'], dtype=object)
```

```
In [53]: from sklearn.metrics import confusion_matrix, accuracy_score
```

```
In [90]: print(confusion_matrix(y_pred, y_test))
```

```
[[934 212]
 [107 156]]
```

```
In [91]: (934+156)/(934+156+212+107)
```

```
Out[91]: 0.7735982966643009
```

```
In [92]: accuracy_score(y_test, y_pred)
```

```
Out[92]: 0.7735982966643009
```

E) Decision Tree:

a. Build a decision tree model where dependent variable is 'Churn' & independent variable is 'tenure'

i. Divide the dataset in 80:20 ratio

ii. Build the model on train set and predict the values on test set

iii. Build the confusion matrix and calculate the accuracy

```
In [93]: x=pd.DataFrame(data.loc[:,['tenure']])
         y=data['Churn']
```

```
In [94]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.8,random_state=0)
```

```
In [95]: from sklearn.tree import DecisionTreeClassifier
```

```
In [96]: # Create Decision Tree classifier object
         clf = DecisionTreeClassifier()

         # Train Decision Tree Classifier
         clf = clf.fit(x_train,y_train)

         #Predict the response for test dataset
         y_pred = clf.predict(x_test)
```

```
In [97]: y_pred
```

```
Out[97]: array(['No', 'No', 'No', ..., 'No', 'No', 'Yes'], dtype=object)
```

```
In [ ]: y_test
```

```
Out[ ]: 2200    No
         4627    No
         3225    No
         2828    No
```

```

3768    No
...
2631    Yes
5333    Yes
6972    Yes
4598    No
3065    No
Name: Churn, Length: 1409, dtype: object

```

```

In [98]: from sklearn.metrics import confusion_matrix, accuracy_score
        from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation

```

```

In [99]: # Model Accuracy, how often is the classifier correct?
        print("Accuracy:", metrics.accuracy_score(y_test, y_pred))

```


Accuracy: 0.7466288147622427

```

In [100]: cm = confusion_matrix(y_test, y_pred)
         plt.figure(figsize=(5,5))
         sns.heatmap(data=cm, linewidths=.5, annot=True, square = True, cmap = 'Blues')
         plt.ylabel('Actual label')
         plt.xlabel('Predicted label')
         all_sample_title = 'Accuracy Score: {0}'.format(clf.score(x_test, y_test))
         plt.title(all_sample_title, size = 15)

```

Out[100]: Text(0.5, 1.0, 'Accuracy Score: 0.7466288147622427')



```

In [125]: !pip install graphviz
         !pip install pydotplus

```

Requirement already satisfied: graphviz in /opt/anaconda3/lib/python3.8/site-packages (0.20)
Requirement already satisfied: pydotplus in /opt/anaconda3/lib/python3.8/site-packages (2.0.2)
Requirement already satisfied: pyparsing>=2.0.1 in /opt/anaconda3/lib/python3.8/site-packages (from pydotplus) (2.4.7)

```

In [126]: pip install scikit-learn==0.20.3

```

ERROR: unknown command "install.scikit-learn==0.20.3"
Note: you may need to restart the kernel to use updated packages.

```

In [129]: from sklearn.tree import export_graphviz
         from sklearn.externals.six import StringIO
         from IPython.display import Image
         import pydotplus

         dot_data = StringIO()
         export_graphviz(clf, out_file=dot_data, filled=True, rounded=True, special_characters=True, feature_names = x_train.columns,
                        class_names=['No', 'Yes'],
                        class_color=[0.5, 0.5, 0.5],
                        node_ids=False)
         graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
         graph.write_png('diabetes.png')

```

```
Image(graph.create_png())
```

```
-----  
ModuleNotFoundError                                Traceback (most recent call last)  
<ipython-input-129-df236d1b4dd6> in <module>  
      1 from sklearn.tree import export_graphviz  
----> 2 from sklearn.externals.six import StringIO  
      3 from IPython.display import Image  
      4 import pydotplus  
      5  
ModuleNotFoundError: No module named 'sklearn.externals.six'
```

F) Random Forest:

a. Build a Random Forest model where dependent variable is 'Churn' & independent variables are 'tenure' and 'MonthlyCharges'

i. Divide the dataset in 70:30 ratio

ii. Build the model on train set and predict the values on test set

iii. Build the confusion matrix and calculate the accuracy


```
In [114]: from sklearn.model_selection import train_test_split  
  
X=data[['tenure']] # Features  
y=data['Churn'] # Labels  
  
# Split dataset into training set and test set  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```


```
In [115]: #Import Random Forest Model  
from sklearn.ensemble import RandomForestClassifier  
  
#Create a Random Forest Classifier  
clf1=RandomForestClassifier(n_estimators=10)  
  
#Train the model using the training sets y_pred=clf.predict(X_test)  
clf1.fit(X_train,y_train)  
  
y_pred=clf1.predict(X_test)  
  
#Import scikit-learn metrics module for accuracy calculation  
from sklearn import metrics  
  
# Model Accuracy, how often is the classifier correct?  
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7501183151916706

```
In [ ]: cm = confusion_matrix(y_test, y_pred)  
plt.figure(figsize=(5,5))  
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')  
plt.ylabel('Actual label')  
plt.xlabel('Predicted label')  
all_sample_title = 'Accuracy Score: {0}'.format(clf1.score(X_test, y_test))  
plt.title(all_sample_title, size = 15)
```

Out[]: Text(0.5, 1.0, 'Accuracy Score: 0.7657359204921912')






In [116..

```
# Try different numbers of n_estimators
estimators = np.arange(10, 150, 10)
accuracy = []

for n in estimators:
    clf1.set_params(n_estimators=n)
    clf1.fit(X_train, y_train)
    y_pred=clf.predict(X_test)
    accuracy.append(metrics.accuracy_score(y_test, y_pred))
plt.title("Effect of n_estimators")
plt.xlabel("n_estimator")
plt.ylabel("accuracy")
plt.plot(estimators, accuracy)
```

Out[116.. [<matplotlib.lines.Line2D at 0x7fcf8d3fbd30>]



In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js